



BUILD YOUR OWN JARVIS

The Free Prompt Pack – by Zubair Trabzada · AI Workshop

6 PROMPTS · ONE EVENING · A TALKING AI SECOND BRAIN

This is the exact method I used to build JARVIS with **Claude Code** – boiled down to a prompt sequence anyone can run. Paste them in order, one at a time, and Claude builds the whole thing for you. No coding required. You just supervise.

What you'll have at the end: a **3D knowledge galaxy** of your own notes · a **voice assistant** that answers from them out loud · the **fly-to-source** camera dive · and a personality with a razor wit that calls you "sir."

Before you start (5 minutes)

- **Claude Code** installed and logged in → claude.com/claude-code
- **Google Chrome** – the mic and voice need it (Safari won't cut it)
- **A folder of markdown notes** (Obsidian vault, exported Notion, anything with .md files). None? Prompt 1 has you covered.
- **An Anthropic API key** for the brain → console.anthropic.com (\$5 of credit is plenty; answers cost fractions of a cent)

⚠ **The one rule: never paste your API key into any chat window.** You'll type it into a config file yourself in Prompt 2. Anything pasted into a chat should be treated as exposed. Trust me on this

one.

Open a terminal, make an empty folder, and start Claude Code inside it:

```
mkdir my-jarvis && cd my-jarvis && claude
```

Run the prompts in order. Let each one finish and **test it** before moving on.

PROMPT 1

The Galaxy

Build me an interactive 3D knowledge galaxy from my markdown notes.

My notes folder is: *[PASTE THE FULL PATH – or replace this line with: "I have no notes – generate 25 realistic sample notes about a small business into ./notes first"]*

Requirements:

- A Python 3 script (build.py, standard library only – no pip installs) that scans every .md file and writes viewer/graph-data.js containing `const GRAPH = {nodes: [...], links: [...]}`. Each node: a label from the filename, its folder as a group, and a ~700-character excerpt of the note text. Link two notes when one mentions the other's title or they share `[[wikilinks]]`.
- A single-page viewer (viewer/index.html) using the 3d-force-graph library from a CDN – no npm, no build tools, no frameworks.
- Make it CINEMATIC: black space background with a starfield, nodes glowing and color-coded by group, slow idle drift. Clicking a node flies the camera to it, highlights its neighbors, and opens a side panel showing that note's excerpt.
- A tiny Python standard-library web server (server.py, port 4700) that serves ONLY the viewer/ folder.
- IMPORTANT: give every node a numeric id equal to its position in the nodes array – later features depend on looking nodes up by index.
- When done: start the server and tell me the URL to open in Chrome.

Checkpoint: you're staring at your own knowledge as a galaxy. Fly around it. This alone is worth a screen recording.

PROMPT 2

The Brain

Now give my galaxy a brain – I want to type a question and get an answer from MY notes.

- Create config.json in the project root (NOT inside viewer/ – it must never be reachable from the browser) with placeholders: {"api_key": "PUT-YOUR-KEY-HERE", "model": "claude-opus-4-8"}. I will paste my API key into that file myself – do not ask me for it.
- Add a POST /chat endpoint to server.py: score every note against my question by keyword overlap (title matches weigh extra), take the top 6, and call the Anthropic Messages API (the model from config.json) with a system prompt that says: answer ONLY from these notes, in 2–3 sentences, and admit it when the notes don't cover it.
- The endpoint returns {"answer": "...", "nodes": [array indexes of the notes used]}.
- Keep a short per-session conversation history server-side so follow-up questions work.
- Add a sleek input bar at the bottom of the viewer that shows the answer on screen.
- The API key must never appear in the browser, the HTML, or any file the server serves.
- When done, restart the server and give me 3 example questions to try based on my actual notes.

Then: paste your API key into config.json, hard-reload Chrome, and interrogate your own brain.

No API key? Tell Claude: *"instead of the Anthropic API, make /chat shell out to claude -p so it runs on my Claude Code subscription."* Slower, but free.

PROMPT 3

The Voice 🗣️

Give it a voice – both directions. Free, built into the browser, no paid services:

- Speak every answer out loud with the Web Speech API (speechSynthesis). Prefer a British English voice if the system has one.
- Add a 🗣️ mic button using webkitSpeechRecognition (Chrome): click, speak, and the transcript goes through the same /chat flow.
- Show a status line while it's listening and thinking ("● listening...", "● thinking...").
- Browsers block audio until the user interacts with the page – make sure speech starts reliably after the first click.
- Test-plan me: list what I should click and say to verify both directions work.

Checkpoint: you're now talking to your notes – and they're talking back.

PROMPT 4

The Magic ✨ – fly-to-source

When Jarvis answers from a note, I want the galaxy to PROVE it:

- Use the nodes array that /chat already returns: fly the camera to the top source node, light it up along with its direct neighbors, and open its side panel – so I can see exactly where the answer came from.
- If the answer drew on 4+ notes, light up the whole cluster instead of flying to one.
- Don't read the whole note aloud – the note is on screen. The spoken answer stays short.

Checkpoint: this is the moment people gasp in demos. Ask it something and watch the camera dive.

PROMPT 5

The Personality 🎩

Give it the personality. Update the /chat system prompt so the assistant:

- Is a dry, impeccably polite British butler with a razor wit. Addresses me as "sir" (occasionally, not every sentence). One genuinely funny line beats three bland ones.
- Answers questions about my notes in ONE witty sentence plus the facts – never recites the note back (it's on screen).
- Handles small talk and jokes WITHOUT dragging the camera around the graph – only fly to nodes when I actually asked about my notes.
- Add a boot greeting: when the page loads, it says "Good evening, sir. [N] notes indexed, all present and accounted for" – with the real note count.

PROMPT 6

Total Recall 📝 – bonus

One more power: let me grow the brain by voice.

- When I say or type something starting with "remember that...", POST it to a new /remember endpoint that writes a real markdown note into a captures/ folder inside my notes directory, with a sensible title from the first few words.
- Then add the new node to the galaxy LIVE – born at its most related node's position, with a brief glow pulse – and fly the camera to it.
- Confirm out loud with one witty line.

Try: "remember that prompt packs make excellent free gifts" – and watch a new star get born.

Troubleshooting (90% of all issues)

SYMPTOM	FIX
Mic button does nothing	Chrome → lock icon in the address bar → allow Microphone. Must be Chrome/Edge.
No sound	Click anywhere on the page once, then ask again (browsers block audio before the first interaction).

Page looks stale after a change	Hard reload: Cmd+Shift+R (Mac) / Ctrl+Shift+R (Windows).
"model not configured"	You didn't paste your key into config.json, or left the placeholder text in.
Answers are generic	Your notes folder path was wrong in Prompt 1 – re-run it with the full absolute path.
Anything else	Paste the exact error into Claude Code and say "fix it." That's the whole method.

Where this stops – and JARVIS begins

What you just built is the skeleton. It's real, it works, and it's yours. The JARVIS I use every day is what happens after **hundreds of hours and millions of tokens** of iteration on top of that skeleton:

- 🎬 The **Iron-Man reactor HUD** that pulses in rhythm with his voice
- 🗣️ A **cloned ElevenLabs voice** + wake word ("Jarvis...") + true barge-in – interrupt him mid-sentence
- ☀️ **"Good morning" briefings from your REAL Gmail + Calendar**, inbox triage, and a "what needs me?" agent inbox
- 🌐 **Live web research by voice** – "Jarvis, research X" → searched, summarized, sources on a card
- 🖐️ **Agent hands** – drafts emails, researches, repurposes content – with a "do it" confirmation gate
- 👁️ **Screen vision** · 🍏 **Mac control by voice** · 🕒 a **Time Machine** ("what was I doing last Tuesday?")
- 🧠 **Model hot-swap** – "Jarvis, try on Grok" – and TARS-style personality dials
- Plus the click-by-click install guide, every update he gets, and the community building alongside you

Jarvis is not for sale – members get him.

The finished build, the 5-minute install, and every future version live inside the **AI Workshop**.

👉 skool.com/aiworkshop

Built with Claude Code + Claude Fable 5. Build your own? Post it in the community – the best one gets featured in the next video. · © Zubair Trabzada